



OpenShift Github Training *Advanced*

Trainer: Alexander Kolin

The World's Local Training Provider

Course objectives

Course objectives

Day 1

- Openshift Architecture
- Network
- Storage
- CRD & Operators
- Kustomize

Course objectives

Day 1

- Openshift Architecture
- Network
- Storage
- CRD & Operators
- Kustomize

Day 2

- Github Actions
- Pod Security
- Storage management
- Advanced Management
- DR & HA Strategies

Course objectives

Day 1

- Openshift Architecture
- Network
- Storage
- CRD & Operators
- Kustomize

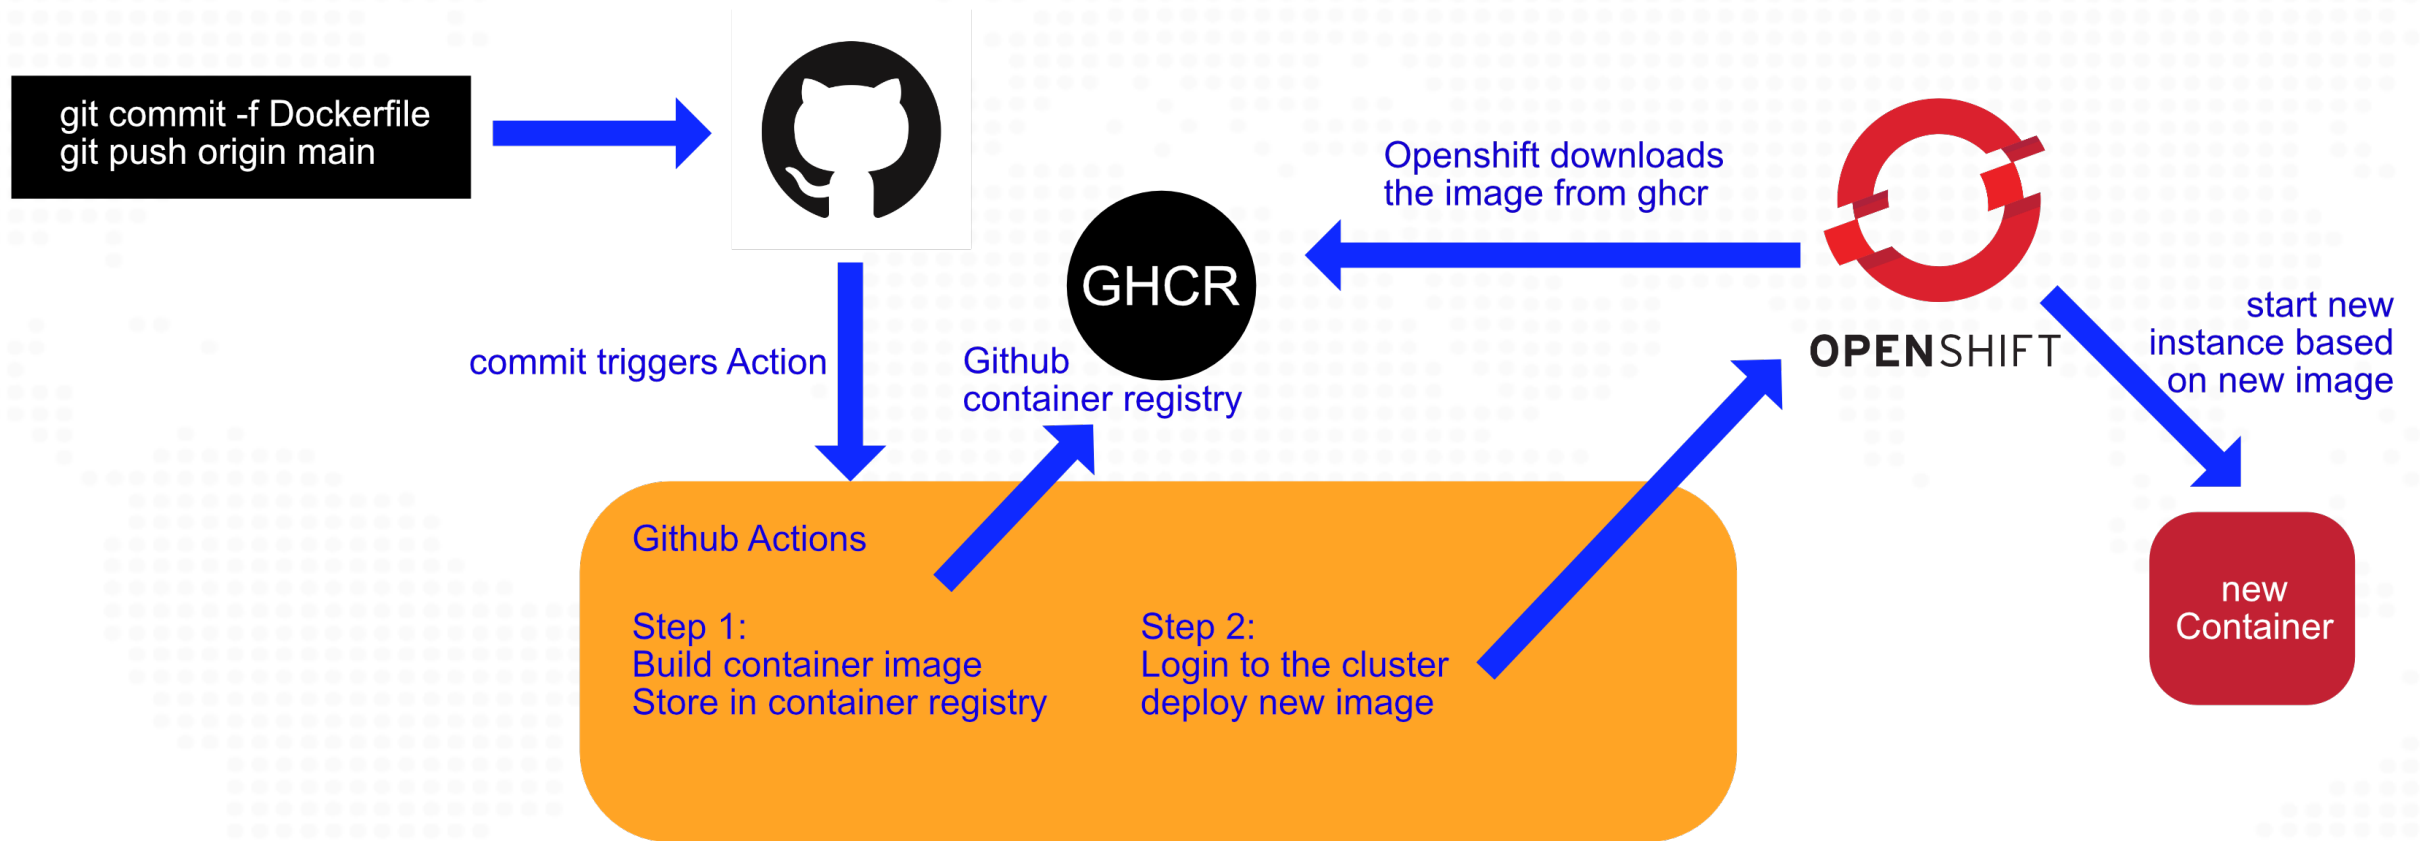
Day 2

- Github Actions
- Pod Security
- Storage management
- Advanced Management
- DR & HA Strategies

Day 3

- App & Infrastructure Security
- Monitoring
- DR & HA Velero
- Wrap Up & Questions

Github



Github Actions Example

```
name: Simple Echo and List Workflow with Git Pull

on:
  push:
    branches:
      - main

jobs:
  echo-and-list-job:
    runs-on: ubuntu-latest
    steps:
      - name: Set up Git
        run: |
          sudo apt-get update
          sudo apt-get install -y git

      - name: Clone Repository using Git
        run: |
          git clone https://github.com/${{ github.repository }}.git
          cd $(basename ${{ github.repository }})
          git pull

      - name: Echo Hello World
        run: echo "Hello, World!"

      - name: List Files
        run: ls *
```

Github Actions Example

```
name: Simple Echo and List Workflow with Git Pull

on:
  push:
    branches:
      - main

jobs:
  echo-and-list-job:
    runs-on: ubuntu-latest
    steps:
      - name: Set up Git
        run: |
          sudo apt-get update
          sudo apt-get install -y git

      - name: Clone Repository using Git
        run: |
          git clone https://github.com/${{ github.repository }}.git
          cd $(basename ${{ github.repository }})
          git pull

      - name: Echo Hello World
        run: echo "Hello, World!"

      - name: List Files
        run: ls *
```

```
name: Simple Echo and List Workflow

on:
  push:
    branches:
      - main

jobs:
  echo-and-list-job:
    runs-on: ubuntu-latest
    steps:
      - name: Checkout Repository
        uses: actions/checkout@v4

      - name: Echo Hello World
        run: echo "Hello, World!"

      - name: List Files
        run: ls *
```


Github Actions Example

```
name: Simple Echo and List Workflow with Git Pull

on:
  push:
    branches:
      - main

jobs:
  echo-and-list-job:
    runs-on: ubuntu-latest
    steps:
      - name: Set up Git
        run: |
          sudo apt-get update
          sudo apt-get install -y git

      - name: Clone Repository using Git
        run: |
          git clone https://github.com/${{ github.repository }}.git
          cd $(basename ${{ github.repository }})
          git pull

      - name: Echo Hello World
        run: echo "Hello, World!"

      - name: List Files
        run: ls *
```

```
name: Simple Echo and List Workflow

on:
  push:
    branches:
      - main

jobs:
  echo-and-list-job:
    runs-on: ubuntu-latest
    steps:
      - name: Checkout Repository
        uses: actions/checkout@v4

      - name: Echo Hello World
        run: echo "Hello, World!"

      - name: List Files
        run: ls *
```

<https://github.com/marketplace>

Actions Example Lab

```
git clone https://github.com/kolinrr/openshift-adv.git
```

```
cd openshift-adv/Day2/0.workflow
```

- Actions -> new Workflow -> set up a workflow
- Copy & paste helloworld.yaml
- Commit

Actions Example Lab

- Actions -> new Workflow -> set up a workflow
- Copy & paste helloworld.yaml
- Commit

```
git clone https://github.com/kolinrr/openshift-adv.git
```

```
cd openshift-adv/Day2/0.workflow
```

```
name: Hello World

on:
  push:
    branches:
      - main

jobs:
  hello:
    runs-on: ubuntu-latest

    steps:
      - name: Checkout Repository
        uses: actions/checkout@v4

      - name: Say Hello
        run: echo "Hello, World!"


      - name: list files
        run: ls -a
```

Github Actions Container Registry

- Github Container Registry -> create Github Packages-Token
 - User-Settings -> Developer -> PAT -> Classic

```
oc create secret docker-registry ghcr-pull-secret \  
--docker-server=ghcr.io \  
--docker-username=GITHUB_USERNAME \  
--docker-password=GITHUB_TOKEN \  
--docker-email=GITHUB_EMAIL
```

- Create Repository Secret:
 - Go to GitHub repository -> Settings -> Secrets and Variables „Actions“
 - GHCR_TOKEN



Github Actions vs Openshift

- Login Openshift <https://developers.redhat.com/developer-sandbox>
- Go to Openshift dashboard, click on your user right on top „copy login command“
 - Copy the data
 - Go to GitHub repository -> Settings -> Secrets and Variables „Actions“
- Create Repository Secrets:
 - OPENSHIFT_SERVER - Copy the string from oc login command at „server“
 - OPENSHIFT_NAMESPACE - its your Openshift project
 - OPENSHIFT_TOKEN - Copy the string at „Your API token is“

Actions Container Image Lab

```
git clone https://github.com/kolinrr/openshift-adv.git
```

```
cd openshift-adv/Day2/0.workflow
```

- Actions -> new Workflow -> set up a workflow
- Copy & paste dockerbasic.yaml
- Set correct Dockerfile path:
e.g. ./Day1/container/Dockerfile
- Commit

Actions Container Image Lab

- Actions -> new Workflow -> set up a workflow
- Copy & paste dockerbasic.yaml
- Set correct Dockerfile path:
e.g. ./Day1/container/Dockerfile
- Commit

```
git clone https://github.com/kolinrr/openshift-adv.git
```

```
cd openshift-adv/Day2/0.workflow
```

```
name: Dockerbasic

env:
...
jobs:

  docker-basic:
    name: Build Dockerimage
    runs-on: ubuntu-latest
    environment: production

    ...
    - name: Build from Dockerfile
      id: build-image
      uses: redhat-actions/buildah-build@v2
      with:
        image: ${ env.APP_NAME }
        tags: ${ env.IMAGE_TAGS }

      dockerfiles: |
        ./Dockerfile
    ...
```

Actions deployment Lab

```
git clone https://github.com/kolinrr/openshift-adv.git
```

```
cd openshift-adv/Day2/0.workflow
```

- Actions -> new Workflow -> set up a workflow
- Copy & paste openshift.yaml
- Commit

Actions deployment Lab

git clone <https://github.com/kolinrr/openshift-adv.git>

cd openshift-adv/Day2/0.workflow

- Actions -> new Workflow -> set up a workflow
- Copy & paste openshift.yaml
- Commit

```
...  
- name: my deployment  
  run: |  
    export IMAGE=${{ steps.push-image.outputs.registry-path }}  
    export OPENSIFT_NAMESPACE=${{ env.OPENSIFT_NAMESPACE }}  
  
    echo ${{ steps.push-image.outputs.registry-path }}  
  
    sed -i 's|IMAGE_PLACEHOLDER|"$IMAGE"|g' Day2/deployment/deployment.yaml  
  
    oc apply -f Day2/deployment/deployment.yaml -n $OPENSIFT_NAMESPACE  
    oc apply -f Day1/network/service.yaml -n $OPENSIFT_NAMESPACE  
  
    oc get pods  
    oc get svc  
...
```

Security ~~PSP~~ vs PSS vs SCC

- PSP - Pod Security Policy - outdated
- Pod Security Standards
 - K8s native
 - Based on Namespace
- SCC
 - Openshift
 - Based on RBAC + Service Account
 - detailed, performant

SCC

Security Context Constraints

- SCC is an **OpenShift-specific** security feature used to control permissions and restrictions for Pods.
- It defines **what a Pod is allowed to do** in terms of security settings.
- Unlike Kubernetes **PodSecurityPolicies (PSP)**, SCCs are **assigned to ServiceAccounts**, not Pods directly.

```
git clone https://github.com/kolinrr/openshift-adv.git
```

```
cd openshift-adv/Day2/1.security
```

```
apiVersion: security.openshift.io/v1
kind: SecurityContextConstraints
metadata:
  name: restricted-scc
allowPrivilegedContainer: false
runAsUser:
  type: MustRunAsNonRoot
seLinuxContext:
  type: MustRunAs
fsGroup:
  type: RunAsAny
supplementalGroups:
  type: RunAsAny
readOnlyRootFilesystem: true
volumes:
  - configMap
  - emptyDir
  - projected
  - persistentVolumeClaim
users:
  - system:serviceaccount:default:secure-sa
```

SCC - Lab

- „Theoretically“ -> On CRC / Sandbox is not possible to test :/

Just go through README.md

```
git clone https://github.com/kolinrr/openshift-adv.git
```

```
cd openshift-adv/Day2/1.security
```

```
apiVersion: security.openshift.io/v1
kind: SecurityContextConstraints
metadata:
  name: restricted-scc
allowPrivilegedContainer: false
runAsUser:
  type: MustRunAsNonRoot
seLinuxContext:
  type: MustRunAs
fsGroup:
  type: RunAsAny
supplementalGroups:
  type: RunAsAny
readOnlyRootFilesystem: true
volumes:
  - configMap
  - emptyDir
  - projected
  - persistentVolumeClaim
users:
  - system:serviceaccount:default:secure-sa
```

Storage Part 2

Performance Testing with FIO Pods

- FIO (**Flexible I/O Tester**) is a benchmarking tool for measuring storage performance.
- Simulates different workloads (random read/write, sequential access, mixed loads).
- Helps identify bottlenecks in storage solutions

Why Use FIO in OpenShift?

- **Test Persistent Volumes (PVs)** before deploying databases or high-IO workloads.
- **Analyze Storage Performance** under real-world conditions.
- **Compare different StorageClasses** to choose the best backend.

Storage Part 2

```
git clone https://github.com/kolinrr/openshift-adv.git
```

```
cd openshift-adv/Day2/2.storage
```

Deployment & benchmark

- See README.md
- `oc exec -it fio-test -- sh`
- `fio --name=randwrite --ioengine=libaio --rw=randwrite --bs=4k --size=1G --numjobs=4 --group_reporting`

Understanding FIO Metrics

IOPS: Number of input/output operations per second

Latency: Response time for read/write operations

Bandwidth: Data transfer speed (MB/s)

CPU Utilization: CPU overhead during storage operations

```
apiVersion: v1
kind: Pod
metadata:
  name: fio-test
spec:
  containers:
    - name: fio-container
      image: nixery.dev/shell/fio
      command: [ "sleep", "3600" ]
      volumeMounts:
        - mountPath: "/data"
          name: storage-volume
  volumes:
    - name: storage-volume
      persistentVolumeClaim:
        claimName: pvc-example
```

Storage Part 2

```
git clone https://github.com/kolinrr/openshift-adv.git
```

```
cd openshift-adv/Day2/2.storage
```

Practical Use Cases

- Evaluate storage performance before database deployments.
- Troubleshoot slow applications due to storage bottlenecks.
- Optimize OpenShift storage configurations.

```
apiVersion: v1
kind: Pod
metadata:
  name: fio-test
spec:
  containers:
    - name: fio-container
      image: nixery.dev/shell/fio
      command: [ "sleep", "3600" ]
      volumeMounts:
        - mountPath: "/data"
          name: storage-volume
  volumes:
    - name: storage-volume
      persistentVolumeClaim:
        claimName: pvc-example
```

Storage Part 2

```
git clone https://github.com/kolinrr/openshift-adv.git
```

```
cd openshift-adv/Day2/2.storage
```

Metrics and Alerts

- Not possible on CRC / Sandbox :/

Environment Management

git clone <https://github.com/kolinrr/openshift-adv.git>

cd openshift-adv/Day2/3.Management

Using Namespaces as separation of the environments

- Better resource isolation
- Access control
- Network security
- Hardware separation

Environment Management

git clone <https://github.com/kolinrr/openshift-adv.git>

cd openshift-adv/Day2/3.Management

MySQL Deployment

- Create namespaces/projects:
 - oc create namespace dev
 - oc create namespace test
 - oc create namespace prod
- Deploy kustomize
 - cd openshift-adv/Day2/3.Management/kustomize
 - oc apply -k overlays/dev ../test ../prod

Environment Management Lab

git clone <https://github.com/kolinrr/openshift-adv.git>

cd openshift-adv/Day2/3.Management

Let's move to README.md

Disaster Recovery High Availability

Velero is

- An open-source tool for backup and restoring Kubernetes cluster resources.
- Supports on-premises and cloud environments.
- Ensures business continuity in case of failures.
- Helps in cluster migrations and rollback strategies.

Disaster Recovery High Availability

What is DR?

- **The ability to restore applications/cluster and data after a failure**
- **Common Causes:**
 - Hardware failures
 - Human errors
 - Security breaches (e.g., ransomware)
 - Natural disasters

Disaster Recovery High Availability

What can Velero do?!

- **Backup entire K8s cluster resources, including volumes data.**
- **Restore application to another cluster**
- **Ensures quick recovery of failed services**
- **Allows cross-region or multi-cluster synch**
- **Provides point-in-time restores**

Disaster Recovery High Availability

Limitations

- Restores can be slow
- Problem with restoring the stateful apps
- Backup of async Apps
- EOF-Error

Disaster Recovery High Availability

Best practices for HA

- Use multi-zone (geo) or multi cluster deployments
- Test the backup process
- Test the restore process
- Monitor backup failures
- Hold „hot“-backup system
- pause the DB while backup



Questions?